Technische Universität Wien

FORSYTE

Formal Methods in Systems Engineering

---

course material: slides

aglaia norza

October 15, 2026

norzaglaia@gmail.com

github.com/AglaiaNorza

Contents

1 Satisfiability in Formal Methods **2**
1.1 The SAT Decision Problem 2
1.2 Formula transformation 2
1.3 Resolution 3
1.4 The DPLL Algorithm 5
1.5 Clause-Driven Conflict Learning 7

2 Algebraic Reasoning **8**

3 Program Analysis **9**

4 Model Checking Approaches **10**

CHAPTER 1

Satisfiability in Formal Methods

1.1 | The SAT Decision Problem

> Def. 1 **The SAT decision problem**
>
> Given a propositional formula in conjunctive normal form, is there a truth assignment to the Boolean variables such that it satisfies all clauses?

SAT is an *NP-complete* problem, and it's a very active research area.

1.2 | Formula transformation

We have that

- every formula can be transformed into an **equivalent** CNF formula with, in the worst case, an exponential increase in size
- every formula can be transformed into an **equisatisfiable** formula with a *linear* increase in size (by introducing new variables for subformulae)

> **Tseitin's transformation**
>
> 1 | introduce a fresh boolean variable for every connective
>
> 2 | encode each equivalence as a set of clauses
>
> 3 | collect all constraints in a big conjunction and enforce the variable representing the entire formula to be true

Example. Consider the following propositional formula:

$$\Phi = (a \wedge (a \vee b)) \vee \neg(a \vee b)$$

First, we break the formula down into an Abstract Syntax Tree (AST) and identify every unique operation, and we then assign a new variable to each (non-atomic) subformula:

$$\begin{aligned} t_0 &\leftrightarrow (t_1 \vee t_3) \\ t_1 &\leftrightarrow (a \wedge t_2) \\ t_2 &\leftrightarrow (a \vee b) \\ t_3 &\leftrightarrow \neg t_4 \\ t_4 &\leftrightarrow (a \vee b) \end{aligned}$$

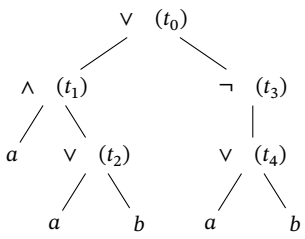

We translate each of these equivalences into standard CNF clauses via:

- AND $X \leftrightarrow (Y \wedge Z)$:
  $\equiv (\neg X \vee Y) \wedge (\neg X \vee Z) \wedge (X \vee \neg Y \vee \neg Z)$
- OR $X \leftrightarrow (Y \vee Z)$:
  $\equiv (\neg X \vee Y \vee Z) \wedge (X \vee \neg Y) \wedge (X \vee \neg Z)$
- NOT $X \leftrightarrow \neg Y$:
  $\equiv (\neg X \vee \neg Y) \wedge (X \vee Y)$

Finally, to assert that the entire formula Φ is true, we assert the root variable (t_0) as a unit clause, and join it with the CNF encodings of all the subformulae:

$$CNF(\Phi) = t_0 \wedge (t_0 \leftrightarrow (t_1 \vee t_3)) \wedge (t_1 \leftrightarrow (a \wedge t_2)) \wedge (t_2 \leftrightarrow (a \vee b)) \wedge (t_3 \leftrightarrow \neg t_4) \wedge (t_4 \leftrightarrow (a \vee b))$$

We get the fully expanded CNF formula:

$$\begin{aligned} CNF(\Phi) = t_0 \wedge\ & \\ & (\neg t_0 \vee t_1 \vee t_3) \wedge (t_0 \vee \neg t_1) \wedge (t_0 \vee \neg t_3) \wedge \\ & (\neg t_1 \vee a) \wedge (\neg t_1 \vee t_2) \wedge (t_1 \vee \neg a \vee \neg t_2) \wedge \\ & (\neg t_2 \vee a \vee b) \wedge (t_2 \vee \neg a) \wedge (t_2 \vee \neg b) \wedge \\ & (\neg t_3 \vee \neg t_4) \wedge (t_3 \vee t_4) \wedge \\ & (\neg t_4 \vee a \vee b) \wedge (t_4 \vee \neg a) \wedge (t_4 \vee \neg b) \end{aligned}$$

This resulting formula is equisatisfiable to the original Φ.

Plasted-Greenbaum's transformation

Every direction in the Tseitin encoding can be replaced by $\rightarrow$ or $\leftarrow$, where the direction of the implication depends on the *polarity* of the formula

- for every occurrence of a subformula φ, the polarity is $+$ if the number of negations above φ in the syntax tree is even, and $-$ otherwise
- we add $t_i \rightarrow$ if the polarity of t_i is positive, and $\leftarrow t_i$ otherwise
- so, every $\rightarrow$ and $\leftarrow$ can be replaced by at most two clauses

1.3 | Resolution

The theory underlying SAT solving is propositional resolution.

Def. 2 **Resolution**

The resolution rule is:

$$\frac{C \cup \{v\} \qquad D \cup \{\neg v\}}{C \cup D}$$

Which tells us that resolving the clause $C \cup \{v\}$ with the clause $D \cup \{\neg v\}$ results in the resolvant clause $D \cup C$.

We also have a **unit-resolution** rule (or **BCP**, boolean constraint propagation). This rule can

be used on so-called *unit clauses*, clauses with a single literal, which needs to be assigned to $\top$ in a CNF.

$$\frac{C \cup \{\neg v\} \qquad \{v\}}{C} \quad \{I, \neg I\} \cap C = \varnothing$$

> NB!!
>
> This rule is *not refutation-complete*, but it can be applied in time *linear* to the size of the formula.

With the unit-resolution rule, we can do **Unit Propagation**: when a unit resolution produces a unit clause, the unit-resolution step can be repeated with the new clause (until fixpoint).

BCP algorithm

- find a unit clause I
- remove all the clauses in which I occurs
- remove all literals $\neg I$ in clauses
- repeat

Example We want to check the validity of:

$$a \wedge b \rightarrow (a \leftrightarrow b)$$

To check validity, we must check unsatisfiability of the negated formula.

1 | CNF of negated formula:

$$\begin{aligned}
&\neg(a \wedge b \rightarrow (a \leftrightarrow b)) \\
\equiv\ & a \wedge b \wedge \neg(a \leftrightarrow b) \\
\equiv\ & a \wedge b \wedge \neg((\neg a \vee b) \wedge (a \vee \neg b)) \\
\equiv\ & a \wedge b \wedge (\neg(\neg a \vee b) \vee \neg(a \vee \neg b)) \\
\equiv\ & a \wedge b \wedge ((a \wedge \neg b) \vee (\neg a \wedge b)) \\
\equiv\ & a \wedge b \wedge (a \vee \neg a) \wedge (a \vee b) \wedge (\neg b \vee \neg a) \wedge (\neg b \vee b) \\
\equiv\ & a \wedge b \wedge (a \vee b) \wedge (\neg a \vee \neg b)
\end{aligned}$$

2 | Unit resolution on a with $(\neg a \vee \neg b)$

$$a \wedge b \wedge (a \vee b) \wedge (\neg a \vee \neg b) \wedge (\neg b)$$

3 | Unit propagation:

$$b \wedge (\neg b)$$

4 | Unit resolution on b with $(\neg b)$

$$\text{empty clause} \bot$$

The negation is UNSAT, therefore the original formula is valid.

1.4 | The DPLL Algorithm

The DPLL (Davis-Putnam-Logemann-Loveland) algorithm followed the DP (Davis-Putnam) algorithm. Where the first eliminated variables via full resolution, causing an explosion in memory, DPLL branches on variable values, trading memory for execution time (DPLL runs in polynomial space, at the cost of a worst case exponential runtime).

DPLL algorithm

- apply repeated unit resolution and removal of subsumed clauses

  at this point, any literals have been removed and the other clauses have been "cleaned up"
- if $\varnothing \in F$, return UNSAT
- if $F = \varnothing$ (no clauses left), return SAT
- choose a new variable x (using a decision heuristic), and branch on it:
  - add $\{x\}$ and recursively solve
  - add $\{\neg x\}$ and recursively solve
- if one of the branches returns SAT, return SAT
- otherwise, return UNSAT

There is an optional optimisation that can be added (but that could be too expensive compared to simple BCP):

Pure Literal Rule

If a variable appears only with positive polarity (or only negative polarity) throughout F:

- setting that literal to $\top$ cannot falsify any clause
- all clauses containing that pure literal can be eliminated immediately

Kinds of assignments in DPLL:

- **Decided ($\bigwedge$):** arbitrary choice/guess when BCP cannot proceed
- **Propagated (|):** values logically forced by unit resolution (BCP) until $\bot$ or a stall

The main drawback of the DPLL algorithm is its **blindness to found conflicts**:

- when a conflict ($\bot$) is encountered, DPLL mechanically undoes only the most recent decision and tries the opposite branch, without analyzing which decisions actually caused the conflict (*chronological backtracking*)
- if a conflict is caused by an independent subset of variables, changing the assignment of an unrelated higher-level variable (like c) forces DPLL to rediscover and prove the exact same conflict all over again
- DPLL does not remember or record why a dead end was reached; it possesses no memory across distinct branches of the search tree

DPLL example: variable choosing Let's consider an example:

$$\mathcal{F} = (\neg b \vee \neg d) \wedge (b \vee \neg d) \wedge (\neg b \vee \neg c) \wedge (b \vee \neg c) \wedge (\neg a \vee c) \wedge (a \vee c \vee d)$$

Initially, no unit clauses exist, so the solver must decide on an unassigned variable.

The choice of an initial variable can heavily influence the search tree:

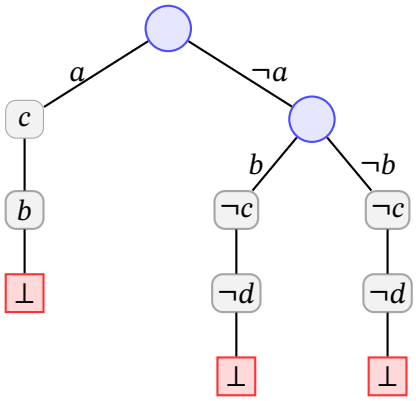

Tree 1: first decision on a
(requires 2 decision levels)

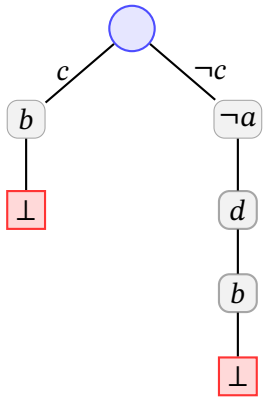

Tree 2: first decision on c
(no secondary decisions needed)

Figure 1.1: DPLL search trees for $\mathcal{F}$. Circles represent decision branch points; rounded rectangles represent literals forced via BCP; and red boxes represent conflict nodes ($\perp$).

DPLL example: blindness to found conflicts Consider the formula:

$$\mathcal{F} = (\neg c \vee a \vee b) \wedge (\neg a \vee \neg b \vee c) \wedge (\neg \mathbf{a} \vee \neg \mathbf{b}) \wedge (\neg \mathbf{a} \vee \mathbf{b}) \wedge (\mathbf{a} \vee \neg \mathbf{b}) \wedge (\mathbf{a} \vee \mathbf{b})$$

Assume the decision heuristic chooses literals in the order: $\neg c \prec c \prec \neg a \prec a \prec \neg b \prec b$.

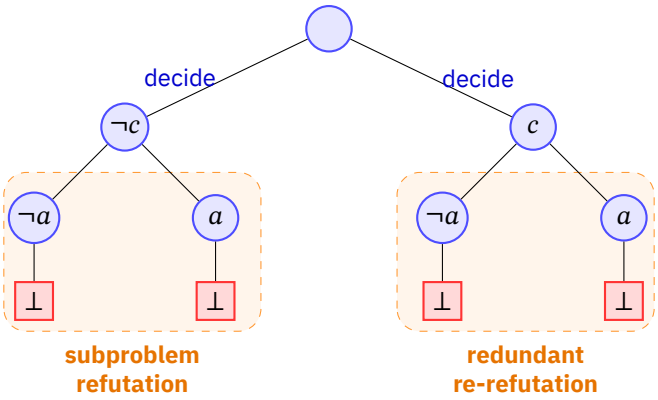

- under $\neg c$, splitting on $\neg a$ and a yields contradictions because the subformula over a, b is unconditionally unsatisfiable
- chronological backtracking flips $\neg c \rightarrow c$.
- cnder c, DPLL has *no memory* of the previous derivation and blindly duplicates the exact same exploration on $\neg a$ and a

1.5 | Clause-Driven Conflict Learning

CHAPTER 2

Algebraic Reasoning

CHAPTER 3

Program Analysis

CHAPTER 4

Model Checking Approaches